

Assignment: Image Classification

- Abigail Smith (ARS190011), Eric Shields (ESS190007)
- CS 4375.004
- Dr. Mazidi
- 04/22/2023

▼ Setting Up Program

```
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sb
import numpy as np
import PIL
import tensorflow as tf
import os

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential

## If training using google colab. Otherwise using pandas
from google.colab import drive
drive.mount('/content/drive')
file = "/content/drive/MyDrive/Colab_Notebooks/tom_and_jerry"

Mounted at /content/drive

## If training from a computer
file = "tom_and_jerry"
```

Describing the Dataset

The dataset used for this assignment is a Tom and Jerry image classification dataset from Kaggle (<https://www.kaggle.com/datasets/balabaskar/tom-and-jerry-image-classification>). The dataset consists of two types of images – either images of Tom or images of Jerry. While the original dataset provided a subset of images with both of these cartoon characters in the same image, we decided for this project to use the subsets with only one of the characters in an image. Both the sequential and CNN models should be able to predict if an image has Tom or has Jerry in it.

▼ Dividing Dataset into Test/Train

```
batchSize = 32 # arbitrary
epochs = 10 # how many passes forward and backwards
imageSize = (200, 100)

train, test = tf.keras.preprocessing.image_dataset_from_directory(
    "/content/drive/MyDrive/Colab_Notebooks/tom_and_jerry",
    validation_split = .2,
    subset = "both",
    seed = 1234,
    image_size = imageSize,
    label_mode = "binary",
    color_mode = "rgb",
    batch_size = batchSize
)

Found 3170 files belonging to 2 classes.
Using 2536 files for training.
Using 634 files for validation.
```

▼ Exploration

```

## Prints the first image batch
for image_batch, labels_batch in train:
    print(image_batch.shape)
    print(labels_batch.shape)
    break

(32, 200, 100, 3)
(32, 1)

## Prints the class names that will be used for training
class_names = train.class_names
print(class_names)

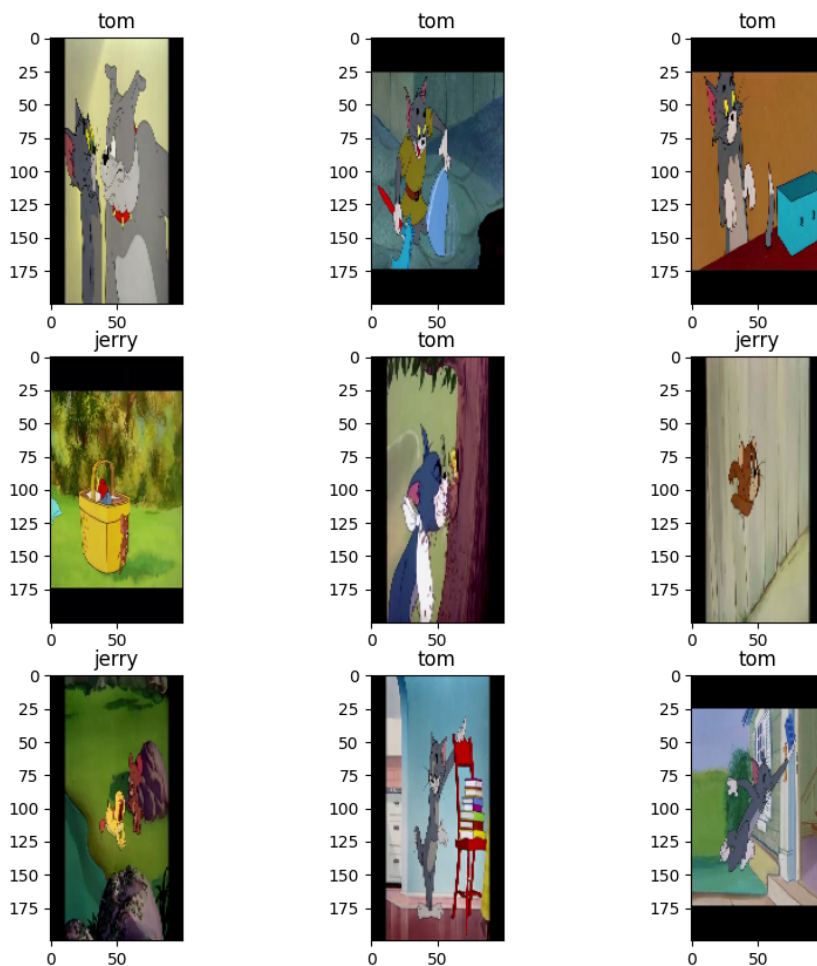
['jerry', 'tom']

```

```

## Prints pictures with all of the various classes
plt.figure(figsize = (10, 10))
for images, labels in train.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[int(labels[i])])

```



```

## Finding distribution in train dataset
labels = []
train_unbatched = tuple(train.unbatch())
for (image, label) in train_unbatched:
    labels.append(label.numpy())
labels = pd.Series(labels)

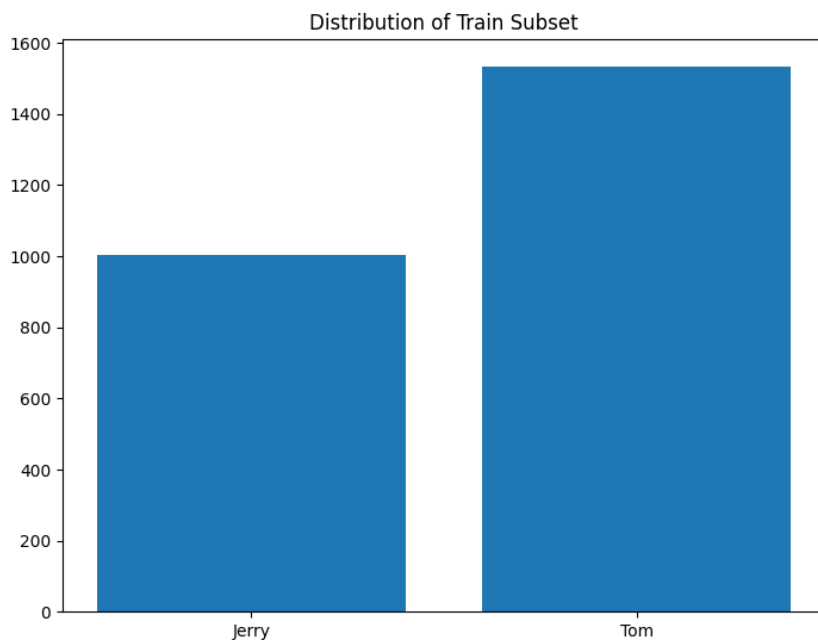
```

```
count = labels.value_counts().sort_index()
count.index = class_names
```

```
print(count)
```

```
jerry    1003
tom      1533
dtype: int64
```

```
figure = plt.figure()
axes = figure.add_axes([1, 1, 1, 1])
x_axis = ['Jerry', 'Tom']
y_axis = [count[0], count[1]]
axes.set_title("Distribution of Train Subset")
axes.bar(x_axis, y_axis)
plt.show()
```

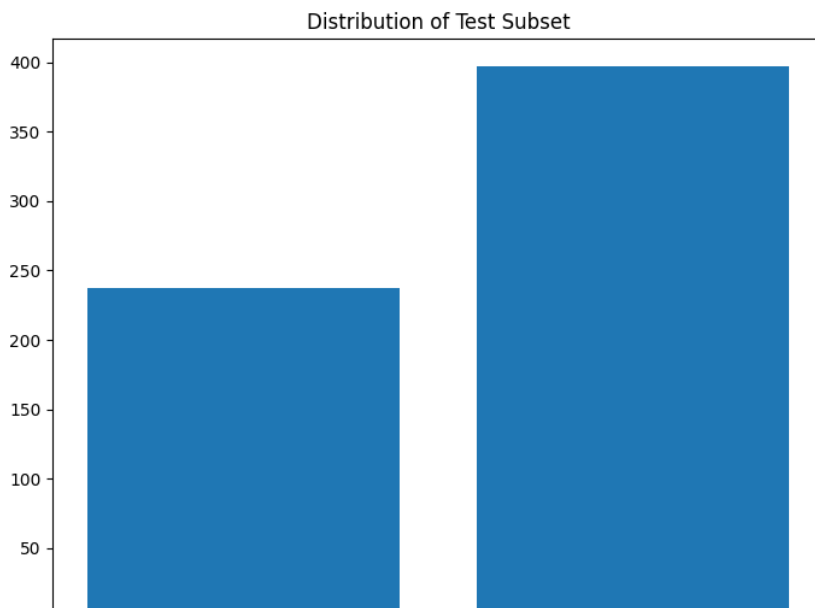


```
## Finding distribution in test dataset
labels = []
test_unbatched = tuple(test.unbatch())
for (image, label) in test_unbatched:
    labels.append(label.numpy())
labels = pd.Series(labels)
count = labels.value_counts().sort_index()
count.index = class_names
```

```
print(count)
```

```
jerry    237
tom      397
dtype: int64
```

```
figure = plt.figure()
axes = figure.add_axes([1, 1, 1, 1])
x_axis = ['Jerry', 'Tom']
y_axis = [count[0], count[1]]
axes.set_title("Distribution of Test Subset")
axes.bar(x_axis, y_axis)
plt.show()
```



▼ Sequential Model

▼ Creating Model

```
#Getting number of classes
num_classes = len(class_names)

## Model creation but without the convolution level for a plain sequential model
model = Sequential([
    layers.Flatten(input_shape = (200, 100, 3)),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(256, activation = 'relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation = 'relu'),
    layers.Dropout(0.2),
    layers.Dense(2, activation = 'softmax')
])
```

▼ Showing Summary

```
## Prints the summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 60000)	0
dense (Dense)	(None, 512)	30720512
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 2)	258

```
=====
Total params: 30,884,994
Trainable params: 30,884,994
Non-trainable params: 0
```

▼ Compiling Model

```
## Compiles the sequential model
model.compile(optimizer='rmsprop',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
```

▼ Fitting Model


```
## Fits the model to the data
epochs = 10
history = model.fit(
    train,
    validation_data=test,
    epochs=epochs,
    batch_size = batchSize,
    verbose = 1
)
```

Epoch 1/10
80/80 [=====] - 56s 673ms/step - loss: 1530.1299 - accuracy: 0.5426 - val_loss: 299.1588 - val_accuracy: 0.3738
Epoch 2/10
80/80 [=====] - 53s 639ms/step - loss: 81.9818 - accuracy: 0.5229 - val_loss: 1.0233 - val_accuracy: 0.3722
Epoch 3/10
80/80 [=====] - 53s 643ms/step - loss: 6.6325 - accuracy: 0.4803 - val_loss: 0.6935 - val_accuracy: 0.3738
Epoch 4/10
80/80 [=====] - 54s 661ms/step - loss: 1.6129 - accuracy: 0.5793 - val_loss: 0.6815 - val_accuracy: 0.6262
Epoch 5/10
80/80 [=====] - 50s 616ms/step - loss: 1.3808 - accuracy: 0.5946 - val_loss: 0.6748 - val_accuracy: 0.6262
Epoch 6/10
80/80 [=====] - 54s 655ms/step - loss: 1.0403 - accuracy: 0.5970 - val_loss: 0.6693 - val_accuracy: 0.6262
Epoch 7/10
80/80 [=====] - 52s 641ms/step - loss: 1.1712 - accuracy: 0.6013 - val_loss: 0.6659 - val_accuracy: 0.6262
Epoch 8/10
80/80 [=====] - 64s 791ms/step - loss: 1.5193 - accuracy: 0.6041 - val_loss: 0.6646 - val_accuracy: 0.6262
Epoch 9/10
80/80 [=====] - 53s 647ms/step - loss: 1.1659 - accuracy: 0.6017 - val_loss: 0.6635 - val_accuracy: 0.6262
Epoch 10/10
80/80 [=====] - 52s 635ms/step - loss: 1.1565 - accuracy: 0.6009 - val_loss: 0.6638 - val_accuracy: 0.6262

▼ Plotting Accuracy to Epoch

```
#Plots the accuracy
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show
```

```
<function matplotlib.pyplot.show(close=None, block=None)>
```



The sequential model is a little different in that the increases we get without the convolution layer are significantly less. Additionally unlike the CNN architecture the model doesn't perform as well on the test data. In fact with a 60 percent success rate we are only slightly more likely to get the right answer than flipping a coin. This was after trying multiple neural net set-ups as well. Tried single Dense layer to 3 layer of Dense with multiple types of nodes at each layer.

▼ CNN Model

▼ Creating Model

```
# Prints the number of classes
num_classes = len(class_names)
```

```
## Makes a sequential model using the keras api for a CNN architecture. Has several layers that ends in a softmax for binary classification
model = Sequential([
    layers.Rescaling(1./255, input_shape=(200, 100, 3)),
    layers.Conv2D(100, (3, 3), activation = "relu"),
    layers.MaxPooling2D(pool_size = (2, 2)),
    layers.Conv2D(100, (3, 3), activation = 'relu'),
    layers.MaxPooling2D(pool_size = (2, 2)),
    layers.Flatten(),
    layers.Dropout(0.5),
    layers.Dense(2, activation = "softmax")
])
```

▼ Showing Summary


```
## The summary of a model
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 200, 100, 3)	0
conv2d (Conv2D)	(None, 198, 98, 100)	2800
max_pooling2d (MaxPooling2D)	(None, 99, 49, 100)	0
conv2d_1 (Conv2D)	(None, 97, 47, 100)	90100
max_pooling2d_1 (MaxPooling2D)	(None, 48, 23, 100)	0
flatten_1 (Flatten)	(None, 110400)	0
dropout_3 (Dropout)	(None, 110400)	0
dense_4 (Dense)	(None, 2)	220802
Total params: 313,702		
Trainable params: 313,702		
Non-trainable params: 0		

▼ Compiling Model

```
## Compiles the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
```



▼ Fitting Model


```
## Fits the model
epochs = 10
history = model.fit(
    train,
    validation_data=test,
    epochs=epochs,
    batch_size = batchSize,
    verbose = 1
)
```

```
Epoch 1/10
80/80 [=====] - 382s 4s/step - loss: 0.6440 - accuracy: 0.6427 - val_loss: 0.4838 - val_accuracy: 0.7823
Epoch 2/10
80/80 [=====] - 333s 4s/step - loss: 0.4826 - accuracy: 0.7780 - val_loss: 0.4874 - val_accuracy: 0.7776
Epoch 3/10
80/80 [=====] - 337s 4s/step - loss: 0.3641 - accuracy: 0.8438 - val_loss: 0.3321 - val_accuracy: 0.8691
Epoch 4/10
80/80 [=====] - 317s 4s/step - loss: 0.2795 - accuracy: 0.8927 - val_loss: 0.2644 - val_accuracy: 0.8896
Epoch 5/10
80/80 [=====] - 344s 4s/step - loss: 0.2055 - accuracy: 0.9263 - val_loss: 0.2279 - val_accuracy: 0.8991
Epoch 6/10
80/80 [=====] - 335s 4s/step - loss: 0.1991 - accuracy: 0.9326 - val_loss: 0.2120 - val_accuracy: 0.9038
Epoch 7/10
80/80 [=====] - 337s 4s/step - loss: 0.1482 - accuracy: 0.9397 - val_loss: 0.2631 - val_accuracy: 0.8975
Epoch 8/10
80/80 [=====] - 315s 4s/step - loss: 0.1350 - accuracy: 0.9507 - val_loss: 0.1932 - val_accuracy: 0.9180
Epoch 9/10
80/80 [=====] - 332s 4s/step - loss: 0.1080 - accuracy: 0.9610 - val_loss: 0.1908 - val_accuracy: 0.9132
Epoch 10/10
80/80 [=====] - 335s 4s/step - loss: 0.0848 - accuracy: 0.9688 - val_loss: 0.1819 - val_accuracy: 0.9211
```

▼ Plotting Accuracy to Epoch

```
## Makes a plot with the accuracy of the model
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show
```

```
<function matplotlib.pyplot.show(close=None, block=None)>
```



When it comes to training this model with the CNN architecture we can see that even the starting accuracy is pretty high at around 70 percent. There are pretty sharp increases as well for the first 2 epochs. After epoch 3 we can see that the model switches to fine tuning the with some

up and down in the train data. What is interesting is that the model definitely isn't overfitted to the dataset, because our test data performs better than the original training dataset.

I think a big reason for this is the convolution layer, but I would be interested to see if changing the image sizes back to their original size would help the model work a little better. Unfortunately I can't test this because I don't have enough ram to keep the pictures in memory.

▼ Using Pretrained Model

by following the provided tutorial: https://www.tensorflow.org/tutorials/images/transfer_learning

▼ Downloading Data

```
# Gets all of the data from a given url and save it as a file
database_url = 'https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip'
path_to_zip = tf.keras.utils.get_file('cats_and_dogs.zip', origin = database_url, extract = True)
PATH = os.path.join(os.path.dirname(path_to_zip), 'cats_and_dogs_filtered')

# Getting train and test directories
train_directory = os.path.join(PATH, 'train')
test_directory = os.path.join(PATH, 'validation')

# Setting up batch and image sizes
batchSize = 32
imageSize = (200, 100)

Downloading data from https://storage.googleapis.com/mledu-datasets/cats\_and\_dogs\_filtered.zip
68606236/68606236 [=====] - 2s 0us/step
```

▼ Creating Train Subset

```
#Loads the data to memory
train = tf.keras.utils.image_dataset_from_directory(
    train_directory,
    shuffle = True,
    batch_size = batchSize,
    image_size = imageSize
)

Found 2000 files belonging to 2 classes.
```

▼ Creating Test Subset

```
#Loads a separate directory as the validation data set
test = tf.keras.utils.image_dataset_from_directory(
    test_directory,
    shuffle = True,
    batch_size = 32,
    image_size = imageSize
)

Found 1000 files belonging to 2 classes.
```

▼ Data Exploration

```
# Gets all of the class names
class_names = train.class_names

## Prints pictures with all of the various classes
plt.figure(figsize = (10, 10))
for images, labels in train.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
```



▼ Preparing to Create Model

```
## Splits the validation dataset and the test set into their batches for training efficiency
test_batches = tf.data.experimental.cardinality(test)
test_dataset = test.take(test_batches // 5)
test = test.skip(test_batches // 5)
```

```
print('Count of test batches: %d' % tf.data.experimental.cardinality(test))
print('Number of test batches: %d' % tf.data.experimental.cardinality(test_dataset))
```

```
Count of test batches: 26
Number of test batches: 6
```


```
## Will prefetch the data
AUTOTUNE = tf.data.AUTOTUNE
```

```
train = train.prefetch(buffer_size = AUTOTUNE)
test = test.prefetch(buffer_size = AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size = AUTOTUNE)
```

```
## Augements the data by spinning the images to a certain extend
data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip('vertical'),
    tf.keras.layers.RandomRotation(0.6),
])
```

```
## Prints an augmented picture
for image, _ in train.take(1):
    plt.figure(figsize = (10, 10))
    first_image = image[0]
    for i in range(9):
```

```
ax = plt.subplot(3, 3, i + 1)
augmented_image = data_augmentation(tf.expand_dims(first_image, 0))
plt.imshow(augmented_image[0] / 255)
```



```
## Makes a preprocessing unit that can be used to scale the images
preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

```
## Rescales the images to be 1/127th the size they were originally
rescale = tf.keras.layers.Rescaling(1./127.5, offset = -1)
```

▼ Creating Model

```
# Creating base model
image_shape = imageSize + (3,)
```

```
base_model = tf.keras.applications.MobileNetV2(
    input_shape = image_shape,
    include_top = False,
    weights = 'imagenet'
)
```

WARNING:tensorflow: `input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224, 3) will be automatically converted to float32 in eager mode. For more details, see https://www.tensorflow.org/api_guides/python/nn_conv2d#input_shape. Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_9406464/9406464 [=====] - 2s 0us/step

```
## Shows that batches are 32 images in a 5 by 5
image_batch, label_batch = next(iter(train))
```

```
feature_batch = base_model(image_batch)
print(feature_batch.shape)
```

```
(32, 7, 4, 1280)
```

▼ Showing Summary

```
# Setting base model to trainable and printing summary
base_model.trainable = False
base_model.summary()
```

```
Model: "mobilenetv2_1.00_224"
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 200, 100, 3)]	0	[]
Conv1 (Conv2D)	(None, 100, 50, 32)	864	['input_1[0][0]']
bn_Conv1 (BatchNormalization)	(None, 100, 50, 32)	128	['Conv1[0][0]']
Conv1_relu (ReLU)	(None, 100, 50, 32)	0	['bn_Conv1[0][0]']
expanded_conv_depthwise (DepthwiseConv2D)	(None, 100, 50, 32)	288	['Conv1_relu[0][0]']
expanded_conv_depthwise_BN (BatchNormalization)	(None, 100, 50, 32)	128	['expanded_conv_depthwise[0][0]']
expanded_conv_depthwise_relu (ReLU)	(None, 100, 50, 32)	0	['expanded_conv_depthwise_BN[0][0]']
expanded_conv_project (Conv2D)	(None, 100, 50, 16)	512	['expanded_conv_depthwise_relu[0][0]']
expanded_conv_project_BN (BatchNormalization)	(None, 100, 50, 16)	64	['expanded_conv_project[0][0]']
block_1_expand (Conv2D)	(None, 100, 50, 96)	1536	['expanded_conv_project_BN[0][0]']
block_1_expand_BN (BatchNormalization)	(None, 100, 50, 96)	384	['block_1_expand[0][0]']
block_1_expand_relu (ReLU)	(None, 100, 50, 96)	0	['block_1_expand_BN[0][0]']
block_1_pad (ZeroPadding2D)	(None, 101, 51, 96)	0	['block_1_expand_relu[0][0]']
block_1_depthwise (DepthwiseConv2D)	(None, 50, 25, 96)	864	['block_1_pad[0][0]']
block_1_depthwise_BN (BatchNormalization)	(None, 50, 25, 96)	384	['block_1_depthwise[0][0]']
block_1_depthwise_relu (ReLU)	(None, 50, 25, 96)	0	['block_1_depthwise_BN[0][0]']
block_1_project (Conv2D)	(None, 50, 25, 24)	2304	['block_1_depthwise_relu[0][0]']
block_1_project_BN (BatchNormalization)	(None, 50, 25, 24)	96	['block_1_project[0][0]']
block_2_expand (Conv2D)	(None, 50, 25, 144)	3456	['block_1_project_BN[0][0]']
block_2_expand_BN (BatchNormalization)	(None, 50, 25, 144)	576	['block_2_expand[0][0]']
block_2_expand_relu (ReLU)	(None, 50, 25, 144)	0	['block_2_expand_BN[0][0]']

```
# Adding a classification head
globalAverageLayer = tf.keras.layers.GlobalAveragePooling2D()
featureBatchAverage = globalAverageLayer(feature_batch)
print(featureBatchAverage.shape) ## Prints the batch sizes
```

```
(32, 1280)
```

```
# Using tf.keras.layers.Dense to change features to a single prediction per image
predictionLayer = tf.keras.layers.Dense(1)
```

Model: "mobilenetv2_1.00_224"

Layer (type) to	Output Shape	Param #	Connected
=====			
input_1 (InputLayer)	(None, 200, 100, 3)	0	[]
Conv1 (Conv2D) ['input_1[0][0]']	(None, 100, 50, 32)	864	
bn_Conv1 (BatchNormalization) ['Conv1[0][0]']	(None, 100, 50, 32)	128	
Conv1_relu (ReLU) ['bn_Conv1[0][0]']	(None, 100, 50, 32)	0	
expanded_conv_depthwise (Depth ['Conv1_relu[0][0]'] wiseConv2D)	(None, 100, 50, 32)	288	
expanded_conv_depthwise_BN (Ba ['expanded_conv_depthwise[0][0]'] tchNormalization)	(None, 100, 50, 32)	128	
expanded_conv_depthwise_relu (['expanded_conv_depthwise_BN[0][0 ReLU)	(None, 100, 50, 32)	0	']
expanded_conv_project (Conv2D) ['expanded_conv_depthwise_relu[0]	(None, 100, 50, 16)	512	[0]']
expanded_conv_project_BN (Batc ['expanded_conv_project[0][0]'] hNormalization)	(None, 100, 50, 16)	64	
block_1_expand (Conv2D) ['expanded_conv_project_BN[0][0]'	(None, 100, 50, 96)	1536]
block_1_expand_BN (BatchNormal ['block_1_expand[0][0]'] ization)	(None, 100, 50, 96)	384	
block_1_expand_relu (ReLU) ['block_1_expand_BN[0][0]']	(None, 100, 50, 96)	0	
block_1_pad (ZeroPadding2D) ['block_1_expand_relu[0][0]']	(None, 101, 51, 96)	0	

block_1_depthwise (DepthwiseCo ['block_1_pad[0][0]'] nv2D)	(None, 50, 25, 96)	864
block_1_depthwise_BN (BatchNor ['block_1_depthwise[0][0]'] malization)	(None, 50, 25, 96)	384
block_1_depthwise_relu (ReLU) ['block_1_depthwise_BN[0][0]']	(None, 50, 25, 96)	0
block_1_project (Conv2D) ['block_1_depthwise_relu[0][0]']	(None, 50, 25, 24)	2304
block_1_project_BN (BatchNorma ['block_1_project[0][0]'] lization)	(None, 50, 25, 24)	96
block_2_expand (Conv2D) ['block_1_project_BN[0][0]']	(None, 50, 25, 144)	3456
block_2_expand_BN (BatchNormal ['block_2_expand[0][0]'] ization)	(None, 50, 25, 144)	576
block_2_expand_relu (ReLU) ['block_2_expand_BN[0][0]']	(None, 50, 25, 144)	0
block_2_depthwise (DepthwiseCo ['block_2_expand_relu[0][0]'] nv2D)	(None, 50, 25, 144)	1296
block_2_depthwise_BN (BatchNor ['block_2_depthwise[0][0]'] malization)	(None, 50, 25, 144)	576
block_2_depthwise_relu (ReLU) ['block_2_depthwise_BN[0][0]']	(None, 50, 25, 144)	0
block_2_project (Conv2D) ['block_2_depthwise_relu[0][0]']	(None, 50, 25, 24)	3456
block_2_project_BN (BatchNorma ['block_2_project[0][0]'] lization)	(None, 50, 25, 24)	96
block_2_add (Add) ['block_1_project_BN[0][0]', 'block_2_project_BN[0][0]']	(None, 50, 25, 24)	0
block_3_expand (Conv2D) ['block_2_add[0][0]']	(None, 50, 25, 144)	3456

```

block_3_expand_BN (BatchNormal (None, 50, 25, 144) 576
['block_3_expand[0][0]']
ization)

block_3_expand_relu (ReLU) (None, 50, 25, 144) 0
['block_3_expand_BN[0][0]']

block_3_pad (ZeroPadding2D) (None, 51, 27, 144) 0
['block_3_expand_relu[0][0]']

block_3_depthwise (DepthwiseCo (None, 25, 13, 144) 1296
['block_3_pad[0][0]']
nv2D)

block_3_depthwise_BN (BatchNor (None, 25, 13, 144) 576
['block_3_depthwise[0][0]']
malization)

block_3_depthwise_relu (ReLU) (None, 25, 13, 144) 0
['block_3_depthwise_BN[0][0]']

block_3_project (Conv2D) (None, 25, 13, 32) 4608
['block_3_depthwise_relu[0][0]']

block_3_project_BN (BatchNorma (None, 25, 13, 32) 128
['block_3_project[0][0]']
lization)

block_4_expand (Conv2D) (None, 25, 13, 192) 6144
['block_3_project_BN[0][0]']

block_4_expand_BN (BatchNormal (None, 25, 13, 192) 768
['block_4_expand[0][0]']
ization)

block_4_expand_relu (ReLU) (None, 25, 13, 192) 0
['block_4_expand_BN[0][0]']

block_4_depthwise (DepthwiseCo (None, 25, 13, 192) 1728
['block_4_expand_relu[0][0]']
nv2D)

block_4_depthwise_BN (BatchNor (None, 25, 13, 192) 768
['block_4_depthwise[0][0]']
malization)

block_4_depthwise_relu (ReLU) (None, 25, 13, 192) 0
['block_4_depthwise_BN[0][0]']

block_4_project (Conv2D) (None, 25, 13, 32) 6144
['block_4_depthwise_relu[0][0]']

block_4_project_BN (BatchNorma (None, 25, 13, 32) 128
['block_4_project[0][0]']

```

```

lization)

block_4_add (Add) (None, 25, 13, 32) 0
['block_3_project_BN[0][0]',
'block_4_project_BN[0][0]']

block_5_expand (Conv2D) (None, 25, 13, 192) 6144
['block_4_add[0][0]']

block_5_expand_BN (BatchNormal (None, 25, 13, 192) 768
['block_5_expand[0][0]']
ization)

block_5_expand_relu (ReLU) (None, 25, 13, 192) 0
['block_5_expand_BN[0][0]']

block_5_depthwise (DepthwiseCo (None, 25, 13, 192) 1728
['block_5_expand_relu[0][0]']
nv2D)

block_5_depthwise_BN (BatchNor (None, 25, 13, 192) 768
['block_5_depthwise[0][0]']
malization)

block_5_depthwise_relu (ReLU) (None, 25, 13, 192) 0
['block_5_depthwise_BN[0][0]']

block_5_project (Conv2D) (None, 25, 13, 32) 6144
['block_5_depthwise_relu[0][0]']

block_5_project_BN (BatchNorma (None, 25, 13, 32) 128
['block_5_project[0][0]']
lization)

block_5_add (Add) (None, 25, 13, 32) 0
['block_4_add[0][0]',
'block_5_project_BN[0][0]']

block_6_expand (Conv2D) (None, 25, 13, 192) 6144
['block_5_add[0][0]']

block_6_expand_BN (BatchNormal (None, 25, 13, 192) 768
['block_6_expand[0][0]']
ization)

block_6_expand_relu (ReLU) (None, 25, 13, 192) 0
['block_6_expand_BN[0][0]']

block_6_pad (ZeroPadding2D) (None, 27, 15, 192) 0
['block_6_expand_relu[0][0]']

```


block_6_depthwise (DepthwiseCo ['block_6_pad[0][0]'] nv2D)	(None, 13, 7, 192)	1728
block_6_depthwise_BN (BatchNor ['block_6_depthwise[0][0]'] malization)	(None, 13, 7, 192)	768
block_6_depthwise_relu (ReLU) ['block_6_depthwise_BN[0][0]']	(None, 13, 7, 192)	0
block_6_project (Conv2D) ['block_6_depthwise_relu[0][0]']	(None, 13, 7, 64)	12288
block_6_project_BN (BatchNorma ['block_6_project[0][0]'] lization)	(None, 13, 7, 64)	256
block_7_expand (Conv2D) ['block_6_project_BN[0][0]']	(None, 13, 7, 384)	24576
block_7_expand_BN (BatchNormal ['block_7_expand[0][0]'] ization)	(None, 13, 7, 384)	1536
block_7_expand_relu (ReLU) ['block_7_expand_BN[0][0]']	(None, 13, 7, 384)	0
block_7_depthwise (DepthwiseCo ['block_7_expand_relu[0][0]'] nv2D)	(None, 13, 7, 384)	3456
block_7_depthwise_BN (BatchNor ['block_7_depthwise[0][0]'] malization)	(None, 13, 7, 384)	1536
block_7_depthwise_relu (ReLU) ['block_7_depthwise_BN[0][0]']	(None, 13, 7, 384)	0
block_7_project (Conv2D) ['block_7_depthwise_relu[0][0]']	(None, 13, 7, 64)	24576
block_7_project_BN (BatchNorma ['block_7_project[0][0]'] lization)	(None, 13, 7, 64)	256
block_7_add (Add) ['block_6_project_BN[0][0]', 'block_7_project_BN[0][0]']	(None, 13, 7, 64)	0
block_8_expand (Conv2D) ['block_7_add[0][0]']	(None, 13, 7, 384)	24576

block_8_expand_BN (BatchNormal ['block_8_expand[0][0]' ization)	(None, 13, 7, 384)	1536
block_8_expand_relu (ReLU) ['block_8_expand_BN[0][0]']	(None, 13, 7, 384)	0
block_8_depthwise (DepthwiseCo ['block_8_expand_relu[0][0]' nv2D)	(None, 13, 7, 384)	3456
block_8_depthwise_BN (BatchNor ['block_8_depthwise[0][0]' malization)	(None, 13, 7, 384)	1536
block_8_depthwise_relu (ReLU) ['block_8_depthwise_BN[0][0]']	(None, 13, 7, 384)	0
block_8_project (Conv2D) ['block_8_depthwise_relu[0][0]']	(None, 13, 7, 64)	24576
block_8_project_BN (BatchNorma ['block_8_project[0][0]' lization)	(None, 13, 7, 64)	256
block_8_add (Add) ['block_7_add[0][0]', 'block_8_project_BN[0][0]']	(None, 13, 7, 64)	0
block_9_expand (Conv2D) ['block_8_add[0][0]']	(None, 13, 7, 384)	24576
block_9_expand_BN (BatchNormal ['block_9_expand[0][0]' ization)	(None, 13, 7, 384)	1536
block_9_expand_relu (ReLU) ['block_9_expand_BN[0][0]']	(None, 13, 7, 384)	0
block_9_depthwise (DepthwiseCo ['block_9_expand_relu[0][0]' nv2D)	(None, 13, 7, 384)	3456
block_9_depthwise_BN (BatchNor ['block_9_depthwise[0][0]' malization)	(None, 13, 7, 384)	1536
block_9_depthwise_relu (ReLU) ['block_9_depthwise_BN[0][0]']	(None, 13, 7, 384)	0
block_9_project (Conv2D) ['block_9_depthwise_relu[0][0]']	(None, 13, 7, 64)	24576

block_9_project_BN (BatchNorma ['block_9_project[0][0]'] lization)	(None, 13, 7, 64)	256
block_9_add (Add) ['block_8_add[0][0]', 'block_9_project_BN[0][0]']	(None, 13, 7, 64)	0
block_10_expand (Conv2D) ['block_9_add[0][0]']	(None, 13, 7, 384)	24576
block_10_expand_BN (BatchNorma ['block_10_expand[0][0]'] lization)	(None, 13, 7, 384)	1536
block_10_expand_relu (ReLU) ['block_10_expand_BN[0][0]']	(None, 13, 7, 384)	0
block_10_depthwise (DepthwiseC ['block_10_expand_relu[0][0]'] onv2D)	(None, 13, 7, 384)	3456
block_10_depthwise_BN (BatchNo ['block_10_depthwise[0][0]'] rmalization)	(None, 13, 7, 384)	1536
block_10_depthwise_relu (ReLU) ['block_10_depthwise_BN[0][0]']	(None, 13, 7, 384)	0
block_10_project (Conv2D) ['block_10_depthwise_relu[0][0]']	(None, 13, 7, 96)	36864
block_10_project_BN (BatchNorm ['block_10_project[0][0]'] alization)	(None, 13, 7, 96)	384
block_11_expand (Conv2D) ['block_10_project_BN[0][0]']	(None, 13, 7, 576)	55296
block_11_expand_BN (BatchNorma ['block_11_expand[0][0]'] lization)	(None, 13, 7, 576)	2304
block_11_expand_relu (ReLU) ['block_11_expand_BN[0][0]']	(None, 13, 7, 576)	0
block_11_depthwise (DepthwiseC ['block_11_expand_relu[0][0]'] onv2D)	(None, 13, 7, 576)	5184
block_11_depthwise_BN (BatchNo ['block_11_depthwise[0][0]'] rmalization)	(None, 13, 7, 576)	2304

block_11_depthwise_relu (ReLU)	(None, 13, 7, 576)	0
['block_11_depthwise_BN[0][0]']		
block_11_project (Conv2D)	(None, 13, 7, 96)	55296
['block_11_depthwise_relu[0][0]']		
block_11_project_BN (BatchNorm	(None, 13, 7, 96)	384
['block_11_project[0][0]']		
alization)		
block_11_add (Add)	(None, 13, 7, 96)	0
['block_10_project_BN[0][0]',		
'block_11_project_BN[0][0]']		
block_12_expand (Conv2D)	(None, 13, 7, 576)	55296
['block_11_add[0][0]']		
block_12_expand_BN (BatchNorma	(None, 13, 7, 576)	2304
['block_12_expand[0][0]']		
lization)		
block_12_expand_relu (ReLU)	(None, 13, 7, 576)	0
['block_12_expand_BN[0][0]']		
block_12_depthwise (DepthwiseC	(None, 13, 7, 576)	5184
['block_12_expand_relu[0][0]']		
onv2D)		
block_12_depthwise_BN (BatchNo	(None, 13, 7, 576)	2304
['block_12_depthwise[0][0]']		
rmalization)		
block_12_depthwise_relu (ReLU)	(None, 13, 7, 576)	0
['block_12_depthwise_BN[0][0]']		
block_12_project (Conv2D)	(None, 13, 7, 96)	55296
['block_12_depthwise_relu[0][0]']		
block_12_project_BN (BatchNorm	(None, 13, 7, 96)	384
['block_12_project[0][0]']		
alization)		
block_12_add (Add)	(None, 13, 7, 96)	0
['block_11_add[0][0]',		
'block_12_project_BN[0][0]']		
block_13_expand (Conv2D)	(None, 13, 7, 576)	55296
['block_12_add[0][0]']		
block_13_expand_BN (BatchNorma	(None, 13, 7, 576)	2304
['block_13_expand[0][0]']		

lization)		
block_13_expand_relu (ReLU)	(None, 13, 7, 576)	0
['block_13_expand_BN[0][0]']		
block_13_pad (ZeroPadding2D)	(None, 15, 9, 576)	0
['block_13_expand_relu[0][0]']		
block_13_depthwise (DepthwiseC	(None, 7, 4, 576)	5184
['block_13_pad[0][0]']		
onv2D)		
block_13_depthwise_BN (BatchNo	(None, 7, 4, 576)	2304
['block_13_depthwise[0][0]']		
rmalization)		
block_13_depthwise_relu (ReLU)	(None, 7, 4, 576)	0
['block_13_depthwise_BN[0][0]']		
block_13_project (Conv2D)	(None, 7, 4, 160)	92160
['block_13_depthwise_relu[0][0]']		
block_13_project_BN (BatchNorm	(None, 7, 4, 160)	640
['block_13_project[0][0]']		
alization)		
block_14_expand (Conv2D)	(None, 7, 4, 960)	153600
['block_13_project_BN[0][0]']		
block_14_expand_BN (BatchNorma	(None, 7, 4, 960)	3840
['block_14_expand[0][0]']		
lization)		
block_14_expand_relu (ReLU)	(None, 7, 4, 960)	0
['block_14_expand_BN[0][0]']		
block_14_depthwise (DepthwiseC	(None, 7, 4, 960)	8640
['block_14_expand_relu[0][0]']		
onv2D)		
block_14_depthwise_BN (BatchNo	(None, 7, 4, 960)	3840
['block_14_depthwise[0][0]']		
rmalization)		
block_14_depthwise_relu (ReLU)	(None, 7, 4, 960)	0
['block_14_depthwise_BN[0][0]']		
block_14_project (Conv2D)	(None, 7, 4, 160)	153600
['block_14_depthwise_relu[0][0]']		
block_14_project_BN (BatchNorm	(None, 7, 4, 160)	640
['block_14_project[0][0]']		
alization)		

block_14_add (Add)	(None, 7, 4, 160)	0
['block_13_project_BN[0][0]',		
'block_14_project_BN[0][0]']		
block_15_expand (Conv2D)	(None, 7, 4, 960)	153600
['block_14_add[0][0]']		
block_15_expand_BN (BatchNorma	(None, 7, 4, 960)	3840
['block_15_expand[0][0]']		
lization)		
block_15_expand_relu (ReLU)	(None, 7, 4, 960)	0
['block_15_expand_BN[0][0]']		
block_15_depthwise (DepthwiseC	(None, 7, 4, 960)	8640
['block_15_expand_relu[0][0]']		
onv2D)		
block_15_depthwise_BN (BatchNo	(None, 7, 4, 960)	3840
['block_15_depthwise[0][0]']		
rmalization)		
block_15_depthwise_relu (ReLU)	(None, 7, 4, 960)	0
['block_15_depthwise_BN[0][0]']		
block_15_project (Conv2D)	(None, 7, 4, 160)	153600
['block_15_depthwise_relu[0][0]']		
block_15_project_BN (BatchNorm	(None, 7, 4, 160)	640
['block_15_project[0][0]']		
alization)		
block_15_add (Add)	(None, 7, 4, 160)	0
['block_14_add[0][0]',		
'block_15_project_BN[0][0]']		
block_16_expand (Conv2D)	(None, 7, 4, 960)	153600
['block_15_add[0][0]']		
block_16_expand_BN (BatchNorma	(None, 7, 4, 960)	3840
['block_16_expand[0][0]']		
lization)		
block_16_expand_relu (ReLU)	(None, 7, 4, 960)	0
['block_16_expand_BN[0][0]']		
block_16_depthwise (DepthwiseC	(None, 7, 4, 960)	8640
['block_16_expand_relu[0][0]']		
onv2D)		
block_16_depthwise_BN (BatchNo	(None, 7, 4, 960)	3840
['block_16_depthwise[0][0]']		

```

rmalization)

block_16_depthwise_relu (ReLU) (None, 7, 4, 960) 0
['block_16_depthwise_BN[0][0]']

block_16_project (Conv2D) (None, 7, 4, 320) 307200
['block_16_depthwise_relu[0][0]']

block_16_project_BN (BatchNorm (None, 7, 4, 320) 1280
['block_16_project[0][0]']
alization)

Conv_1 (Conv2D) (None, 7, 4, 1280) 409600
['block_16_project_BN[0][0]']

Conv_1_bn (BatchNormalization) (None, 7, 4, 1280) 5120
['Conv_1[0][0]']

out_relu (ReLU) (None, 7, 4, 1280) 0
['Conv_1_bn[0][0]']

=====
=====
Total params: 2,257,984
Trainable params: 0
Non-trainable params: 2,257,984

```

```
predictionBatch = predictionLayer(featureBatchAverage)
print(predictionBatch.shape)
```

```
(32, 1)
```

▼ Compiling Model

```
# Building model
inputs = tf.keras.Input(shape=(200, 100, 3))
x = data_augmentation(inputs)
x = preprocess_input(x)
x = base_model(x, training=False)
x = globalAverageLayer(x)
x = tf.keras.layers.Dropout(0.2)(x)
outputs = predictionLayer(x)
model = tf.keras.Model(inputs, outputs)

# Compiling model
learningRate = 0.0001
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=learningRate),
    loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
    metrics=['accuracy'])
```

```
# Printing summary
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[None, 200, 100, 3]	0
sequential_1 (Sequential)	(None, 200, 100, 3)	0
tf.math.truediv (TFOpLambda)	(None, 200, 100, 3)	0
tf.math.subtract (TFOpLambda)	(None, 200, 100, 3)	0
mobilenetv2_1.00_224 (Functional)	(None, 7, 4, 1280)	2257984
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
dropout_1 (Dropout)	(None, 1280)	0
dense_1 (Dense)	(None, 1)	1281

```
=====  
Total params: 2,259,265  
Trainable params: 1,281  
Non-trainable params: 2,257,984  
=====
```

▼ Fitting Model

```
## Fits the model
epochs = 10
history = model.fit(train,
                    epochs=epochs,
                    validation_data=test)

Epoch 1/10
63/63 [=====] - 58s 838ms/step - loss: 0.9881 - accuracy: 0.4700 - val_loss: 0.8405 - val_accuracy: 0.4678
Epoch 2/10
63/63 [=====] - 59s 920ms/step - loss: 0.8186 - accuracy: 0.5380 - val_loss: 0.6594 - val_accuracy: 0.6052
Epoch 3/10
63/63 [=====] - 59s 931ms/step - loss: 0.7639 - accuracy: 0.5705 - val_loss: 0.5421 - val_accuracy: 0.6832
Epoch 4/10
63/63 [=====] - 60s 951ms/step - loss: 0.6733 - accuracy: 0.6195 - val_loss: 0.4619 - val_accuracy: 0.7537
Epoch 5/10
63/63 [=====] - 61s 970ms/step - loss: 0.6346 - accuracy: 0.6535 - val_loss: 0.4221 - val_accuracy: 0.7686
```



```

Epoch 6/10
63/63 [=====] - 59s 941ms/step - loss: 0.6164 - accuracy: 0.6660 - val_loss: 0.3752 - val_accuracy: 0.8045
Epoch 7/10
63/63 [=====] - 58s 926ms/step - loss: 0.5695 - accuracy: 0.6940 - val_loss: 0.3335 - val_accuracy: 0.8354
Epoch 8/10
63/63 [=====] - 59s 941ms/step - loss: 0.5637 - accuracy: 0.6960 - val_loss: 0.3210 - val_accuracy: 0.8329
Epoch 9/10
63/63 [=====] - 60s 947ms/step - loss: 0.5379 - accuracy: 0.7235 - val_loss: 0.3058 - val_accuracy: 0.8403
Epoch 10/10
63/63 [=====] - 53s 850ms/step - loss: 0.5270 - accuracy: 0.7360 - val_loss: 0.2950 - val_accuracy: 0.8416

```


▼ Plotting Accuracy and Loss to Epoch

```

## Makes a plot with the accuracy of the model
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show

```

```
prev_acc = history.history['accuracy']
```




```

## Makes a plot with the loss of the model
plt.plot(history.history['val_loss'])
plt.plot(history.history['loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show

```

```
prev_loss = history.history['loss']
```



This model uses the CNN model previously constructed. We can see that at least in regards to its current data of cats and dogs it has a success rate of around 75 percent after 10 epochs. We can see from the loss diagram that, as the model progressed through the epochs, that the model was able to improve in its predictions – a fact that is reflected in the way accuracy rised throughout each epoch. Though not required, in the next steps we are going to fine tune the model to attempt to see if we can improve the performance of the model at all.

Additional Steps: Fine Tuning

Preparing for Fine Tuning

```
## Allows for finetuning from other people or ourselves
base_model.trainable = True

# Unfreezing the base model and setting bottom layers to untrainable
fine_tune_at = 100
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False
```

Compiling Model

```
model.compile(
    loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
    optimizer = tf.keras.optimizers.RMSprop(learning_rate=learningRate/10),
    metrics=['accuracy']
)
```

Showing Summary


```
model.summary()

## Model has 8 different levels to train and most are transformation layers though and
## only at the mobilenet and dense layers are we training the model
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 200, 100, 3)]	0
sequential_1 (Sequential)	(None, 200, 100, 3)	0
tf.math.truediv (TFOpLambda)	(None, 200, 100, 3)	0
tf.math.subtract (TFOpLambda)	(None, 200, 100, 3)	0
mobilenetv2_1.00_224 (Functional)	(None, 7, 4, 1280)	2257984
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
dropout_1 (Dropout)	(None, 1280)	0
dense_1 (Dense)	(None, 1)	1281

Total params: 2,259,265



Trainable params: 1,862,721
 Non-trainable params: 396,544

▼ Fitting Model

We have to pass through the hyper parameters that will define the new model. This shows through as additional epochs 11-20
 ## that can be seen as additional training

```
tunedEpochs = 10
totalEpochs = epochs + tunedEpochs
```

```
history_fine = model.fit(
    train,
    epochs=totalEpochs,
    initial_epoch=history.epoch[-1],
    validation_data=test
)
```


```
Epoch 10/20
63/63 [=====] - 99s 1s/step - loss: 0.4866 - accuracy: 0.7690 - val_loss: 0.2118 - val_accuracy: 0.8936
Epoch 11/20
63/63 [=====] - 92s 1s/step - loss: 0.3859 - accuracy: 0.8125 - val_loss: 0.1547 - val_accuracy: 0.9356
Epoch 12/20
63/63 [=====] - 86s 1s/step - loss: 0.3524 - accuracy: 0.8360 - val_loss: 0.1558 - val_accuracy: 0.9381
Epoch 13/20
63/63 [=====] - 89s 1s/step - loss: 0.3420 - accuracy: 0.8395 - val_loss: 0.1635 - val_accuracy: 0.9134
Epoch 14/20
63/63 [=====] - 92s 1s/step - loss: 0.3268 - accuracy: 0.8455 - val_loss: 0.1499 - val_accuracy: 0.9369
Epoch 15/20
63/63 [=====] - 92s 1s/step - loss: 0.2991 - accuracy: 0.8645 - val_loss: 0.1442 - val_accuracy: 0.9307
Epoch 16/20
63/63 [=====] - 91s 1s/step - loss: 0.3091 - accuracy: 0.8525 - val_loss: 0.1553 - val_accuracy: 0.9455
Epoch 17/20
63/63 [=====] - 99s 2s/step - loss: 0.3062 - accuracy: 0.8635 - val_loss: 0.1470 - val_accuracy: 0.9394
Epoch 18/20
63/63 [=====] - 85s 1s/step - loss: 0.2901 - accuracy: 0.8745 - val_loss: 0.1439 - val_accuracy: 0.9455
Epoch 19/20
63/63 [=====] - 86s 1s/step - loss: 0.2794 - accuracy: 0.8740 - val_loss: 0.1333 - val_accuracy: 0.9443
Epoch 20/20
63/63 [=====] - 92s 1s/step - loss: 0.2611 - accuracy: 0.8820 - val_loss: 0.1310 - val_accuracy: 0.9455
```

▼ Plotting Accuracy and Loss to Epoch

```
## Makes a plot with the accuracy of the model
plt.plot(history_fine.history['val_accuracy'])
plt.plot(history_fine.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show
```

```
<function matplotlib.pyplot.show(close=None, block=None)>
```


Model Accuracy



```
## Makes a plot with the loss of the model
plt.plot(history_fine.history['val_loss'])
plt.plot(history_fine.history['loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc = 'upper left')
plt.show
```

```
<function matplotlib.pyplot.show(close=None, block=None)>
```

Model Loss



With fine tuning, we were able to steadily increase the accuracy of our model and, in turn, decrease the loss. One important thing to note is that the test subset has the greatest response to fine tuning, as seen in both the "Model Accuracy" and "Model Loss" graphs. In contrast, the train subset had little impact, but steadily showed improvement with each epoch.

✓ 2s completed at 1:48 PM



